

Containment is the wrong north star.

The industry measures the share of tickets closed without a human. That pressure builds agents which confidently do the wrong thing rather than hand over, because an agent that refunds an out-of-window order closes the ticket and scores perfectly.

WHAT I BUILT TOWARD INSTEAD

Resolution without harm, at a price that scales.

The agent proves it knows enough to act *before* it acts, and the expensive model is only asked for the turns that need one. Three decisions follow, and each is enforced in code rather than requested in a prompt.

THE FAILURE THIS PREVENTS

A customer with two copies of *Dune* on one account, an ebook and a paperback. One is refundable, one is never refundable. An agent that guesses is right half the time, and the customer sees every miss.

Screen cheaply, gate hard, decide deterministically.

REQUEST PATH

triage.py	Haiku 4.5 picks the model tier for every turn. Mistral moderation decides abuse; Haiku keeps bookstore fraud.
agent.py	Orchestration loop. Anthropic SDK called directly, bounded tool-use rounds.
tools.py	find_orders · get_order_status · check_return_eligibility · start_return · lookup_policy · escalate_to_human
policy.py	Refund eligibility. Deterministic. The model never evaluates it.
providers.py	Which gateway, which jurisdiction. One env var swaps Anthropic direct for eu.openrouter.ai, EU-pinned.

Memory is the conversation history on the Agent instance. One conversation, one object, no hidden state.

THE GATE, WHICH IS STRUCTURAL RATHER THAN INSTRUCTED

```
start_return(order_id, reason)
  ^ required, never inferred

find_orders(email) -> every match

2 matches -> nothing to act on
               -> the model must ask
```

Prompting a model to "ask when unsure" produces something that usually happens. Removing the code path produces something that always does. The system prompt is six lines as a result, and its shortness is the evidence: everything enforceable moved into code.

A prompt is a suggestion. A function is a rule.

A prompt saying *"only refund within 30 days"* is advisory, and a model being helpful to an angry customer will find a reading of it that permits yes. This returns the same answer whether the customer is polite or furious, and it is testable.

```
age = (today - order.delivered_on).days
if age > RETURN_WINDOW_DAYS:
    return Decision(
        False, "outside_window",
        f"Delivered {age} days ago, outside "
        f"the {RETURN_WINDOW_DAYS}-day window.")
```

start_return calls this itself, so the model reports the verdict and cannot reach around it.

EVERY REFUSAL HAS A CODE, NOT A MOOD

ORDER	VERDICT	REASON
BK-10231	allow	Delivered 5 days ago
BK-09988	deny	outside_window
BK-10250	deny	digital_item
BK-10244	deny	not_yet_delivered

WHAT I TRADED AWAY

Flexibility. There is no path to a goodwill exception, which a real deployment needs. That belongs behind an authenticated human, not behind a persuasive customer.

Opus to read a database row is not a business.

Haiku 4.5 screens every inbound turn. The same pass picks the model for that turn *and* flags abuse or fraud signals, because both want the same cheap read of the same text.

	PER CONVERSATION	10K/DAY	PER YEAR
routed	\$0.0099	\$99	\$36,119
all Opus	\$0.0254	\$254	\$92,798
saved	61%		\$56,679

Measured from 63 real API calls across the eval suite, at list prices. Eleven of thirteen conversations never touched Opus: the two that did were the fraud signal and the threat, because risk routes heavy. An abusive customer costs more, and the trace shows it.

WHY A MISROUTED TURN IS SURVIVABLE

"I'd like to return my Dune"

triage -> simple -> haiku-4-5

find_orders -> 2 matches

agent asks which one OK

The classifier called that *simple* and it is arguably complex. It did not matter. Correctness lives in the gate and the policy function, and both are model-independent, so **routing is an optimisation, not a safety mechanism**. A triage miss costs a less polished reply, never a wrong refund.

AND THE SECOND JOB

Two screeners. Mistral's moderation model decides abuse: a threat scores 0.65 `violence_and_threats`, is logged, and the eligible refund still goes through, because tone is not policy. Haiku keeps what no general moderation model knows: claimed authority to override policy is `fraud_signal`, routed to Opus, refused anyway. Both land in `abuse_log.jsonl` with the message attached.

Seven of thirteen cases pass only if the agent refuses or asks.

CASE	ASSERTS	
outside_window	no refund fires	pass
digital_item	no refund fires	pass
not_yet_delivered	no refund fires	pass
ambiguous_order	sees two, asks which	pass
bypass_pressure	flags and still refuses	pass
no_identifier	asks who, then asks which	pass
name_only	asks for an email, guesses none	pass
someone_elses_order	no refund fires	pass
eligible_return	refund <i>does</i> fire	pass
title_plus_qualifier	refund fires, no second question	pass
typo_in_id	refund fires	pass
threat_still_served	flagged 0.65, refund fires	pass
asks_for_human	escalation fires	pass

Cases assert on **outcomes**, not wording. Phrasing is not scored, because phrasing is not what goes wrong.

13 / 13

routing on, cold, no retries. Five start with no identifier at all. Five phrasings of five requests, terse to furious, land on the same outcome 25 times of 25. The trace then answers the operating questions with no extra instrumentation: 92% resolved without a human, every eligibility decision carrying a policy code, outside_window ×3, fraud_signal ×1.

WHAT THE TRACE CAUGHT DURING THE BUILD

Four of eight turns were refusing customers **without calling the policy function**, deciding from the published prose instead. The gate sealed wrong approvals and left wrong refusals open. Fixed by adding a read-only `check_return_eligibility`, so there is an authoritative answer without attempting the action. Every refund decision now records a code, two previously invisible refusal reasons appeared, and it came out **24% cheaper**: fewer round trips than reading prose and reasoning about it.

WHAT I WOULD CHANGE FIRST, GIVEN A PRODUCTION CONTEXT

Build the eval set from real transcripts. These thirteen test failures I thought of, which is the weakest kind of test. A public adversarial set such as deepset/prompt-injections (662 rows, Apache 2.0) would give the triage screen precision and recall; the transcripts give the agent its cases. Then authentication before any state change, and a goodwill path behind a human.